



InvestPuppy
investpuppy.com

UNVARNISHED

The invisible stakeholders

Serious when it matters.

About InvestPuppy

InvestPuppy builds Vektor, a systematic portfolio management platform for wealth management professionals.

What this series is

Every experienced practitioner thinks it. Almost nobody publishes it: *there has to be a better way.*

We thought it long enough that we went and built one.

Before that, we spent decades in the rooms.

The same project was failing for the third time. The project plan had been signed off before anyone had understood the first requirement.

The relationship was too warm for anyone to say so.

It isn't incompetence.

It's the incentive structure doing its job perfectly.

Nobody planned it this way. Nobody needed to.

We got frustrated enough to write it down. Then frustrated enough to build something. These papers are the writing-it-down part.

The something we built is a separate conversation.

Every financial software implementation has two sets of stakeholders. The first set attended the workshops, reviewed the specification, and signed the project plan. The second set did not attend anything. They were not invited, or were invited once and sent a delegate, or were told the project was being handled by the first set and they would be consulted when relevant. They discovered the system on go-live day. They have been explaining why it does not work for them ever since.

1. Who Is Not in the Room

The workshop that produces the specification is not a representative sample of the people who will use the system. It is a sample of the people who were available, willing, and sufficiently senior to attend. These are not the same people.

The operations director attends. The portfolio analyst who processes thirty transactions a day does not. The head of compliance attends. The compliance officer who produces the monthly regulatory report does not. The CTO attends the architecture session. The system administrator who will

maintain the integrations does not.

The people who attend the workshops describe the system at the level at which they interact with it – which is to say, the level at which they rarely interact with it at all. The people who interact with it daily, in detail, at volume, were not asked.

The specification that results from a workshop attended by senior stakeholders will correctly describe the high-level workflow. It will incompletely describe the exceptions, the edge cases, the manual processes, and the reporting requirements that matter most to the people who use the system every day. These are the things that generate post-go-live complaints. They were always present. Nobody asked.

2. The Map Is Not the Stakeholder

The stakeholder map is not the same as the stakeholder. This distinction is where most stakeholder engagement processes fail.

Stakeholder mapping produces a document. The document lists the stakeholders, their level of interest, their level of influence, and their communication plan. It is produced in week two of the project and updated in week three. It is a governance artefact. It is not an engagement process.

The compliance function listed on the stakeholder map as ‘high interest / medium influence’ with a bi-weekly update did not attend the configuration workshops. They received the updates. They may have read some of them. They were not asked whether the system would produce the regulatory report they need. The go-live post-implementation review will note that the compliance reporting requirement was not in scope, was raised as a post-go-live item, and is currently being assessed for inclusion in phase two.

Listing a stakeholder on a map, sending them a fortnightly summary, and asking them to confirm receipt is not stakeholder engagement. It is stakeholder administration. The two things are not the same, and confusing them is one of the primary reasons invisible stakeholders remain invisible.

3. The Empty Chair

There is a moment in every failed implementation where the invisible stakeholder becomes visible. It is usually about three weeks after go-live. The system is live. The immediate crisis of the cutover has passed. The project team is beginning to wind down. And then someone – a compliance officer, a back-office analyst, a reporting team lead – submits a ticket, or attends a post-go-live meeting, or walks into the project manager’s office and says: this system does not do the thing I need it to do.

They are not wrong. The system does not do the thing they need it to do, because they were not in the room when the configuration decisions were made.

There is a more expensive version of this problem. It does not arrive at go-live. It arrives when regulatory requirements change after the system is live – and the architecture cannot accommodate the change because the compliance function was not involved in the design.

By this point the project is closed, the team is disbanded, and the cost of the remediation is a new project.

The chair had their name on it. It was empty when the room was designed.

4. Finding the Invisible Stakeholders

The invisible stakeholder is not hard to find. They are the person who currently does, manually or in a legacy system, the thing the new system is supposed to replace. They exist in every function the system touches. They are identifiable before the specification is written.

The question that finds them is not 'who is responsible for this function?' It is 'who does this work?' These are frequently different people, and only one of them knows what the system needs to do.

Three things that surface invisible stakeholders before go-live:

One. A workflow walkthrough conducted with the people who do the work, not the people who manage it — specifically including the exception-handling processes that senior stakeholders may not know exist.

Two. A reporting requirement sign-off process that includes the recipients of every report the system will produce, not only the commissioners. The person who uses the report is the authority on whether it is correct.

Three. A UAT process that requires sign-off from every function that will use the system in production. Sign-off by a project team member on behalf of a function is not sign-off. It is documentation of the gap that will become visible at go-live.

Ask who is not in the room. Then go and find them.

The empty chair was always there.

It had a name on it.